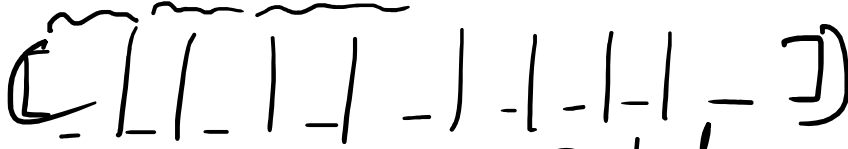


Problems w/ LSTM & RNN

→ Slow

→ Requires a lot of data

→ Accuracy (Vanishing Gradient Problem)

→ 
Tokenize then Embed

→ Transformer: Picks up entire sequence.

→ $\begin{cases} \text{Query: What is the text} \\ \text{Key: Where is the text} \\ \text{Value: Here is the text} \end{cases}$ (Systematic Syntax & Derivable Semantics)

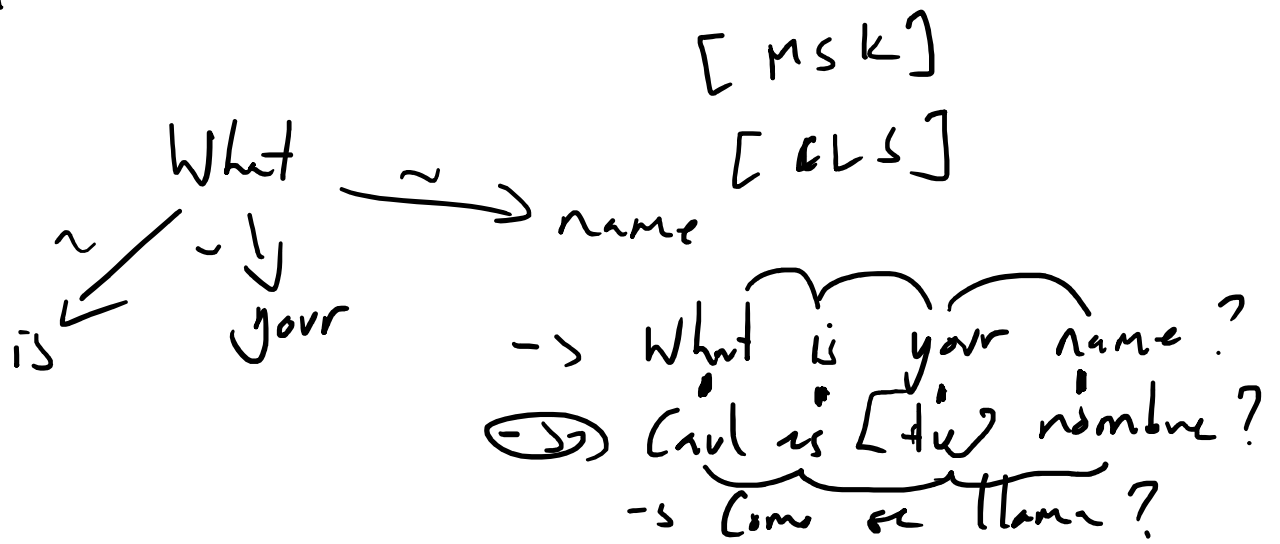
→ Encoder & Decoder

Q, K, V tries to understand Q, K, V.

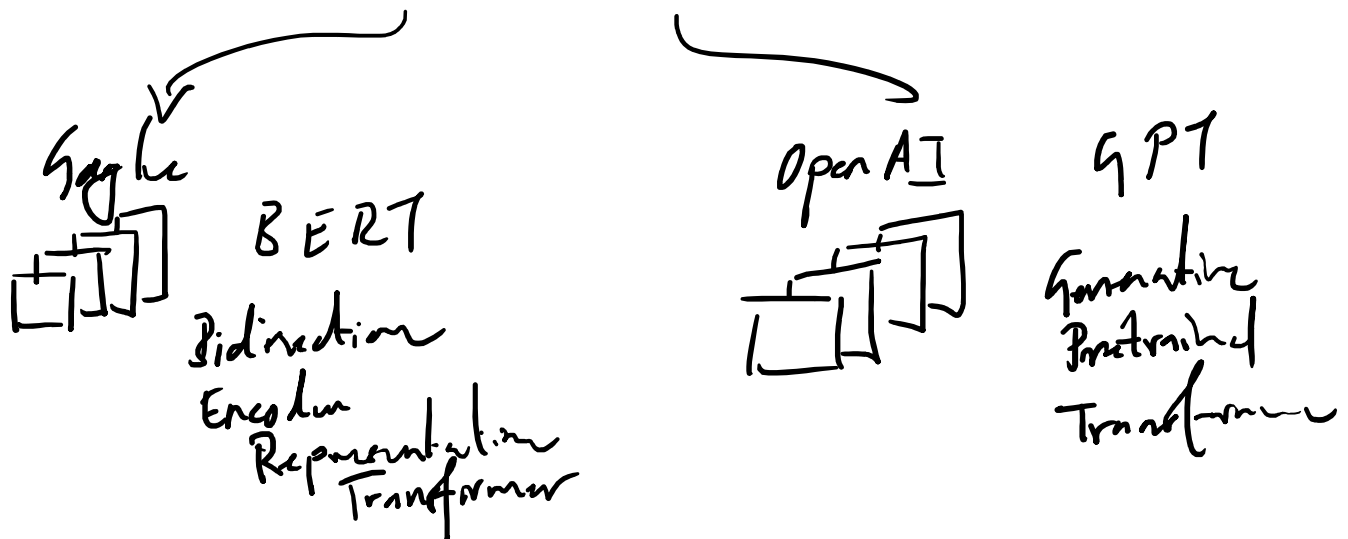
→ GAN ~ Generative Adversarial Networks.

→ Masking What is your name?

Multi-Headed Attention



→ Separate Encoders & Decoders



→ Machine Translation & Text-generation (GPT)

→ Text Classification & Text Summarization (BERT)

→ GPT-1, GPT-2, GPT-3, GPT-4 ~ Not source.

→ BERT, RoBERTa, XLNet, DistilBERT ~ Hugging Face

→ BERT, RoBERTa, XLNet, DistilBERT ~ Fluggly tree
110M 135M 102h GPU ~ V100

→ GPT-3.5 & GPT-4 ~ FAIR ~ Yann LeCun

→ Meta - LLaMa - 3.5 B parameters
Wikipedia & News ~ 80 B tokens
V100 & T100

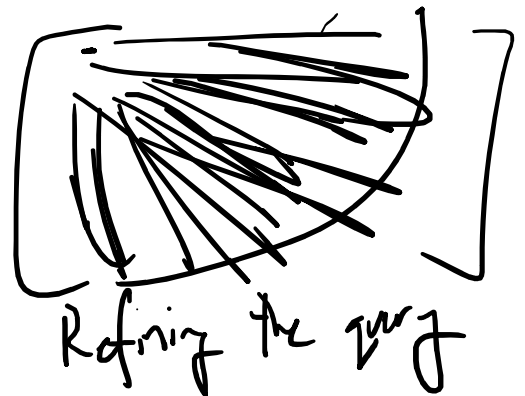
→ LLM ~ Large Language Model (High tech Transformer)
~ Manifests - Use & calibrate this model.

[Agri / Automobiles / Electronics]
[7000000000]

~ Tell me about Lamb.

→ Prompt Engineering

[GPT ~ OpenAI - Paid
Claude ~ Anthropic ~ Paid
✓ LLaMa ~ Meta ~ Open Source
Gemma ~ Google ~ OS
Vistra AI ~ Paid]



→ Llama Index
→ Langchain ✓

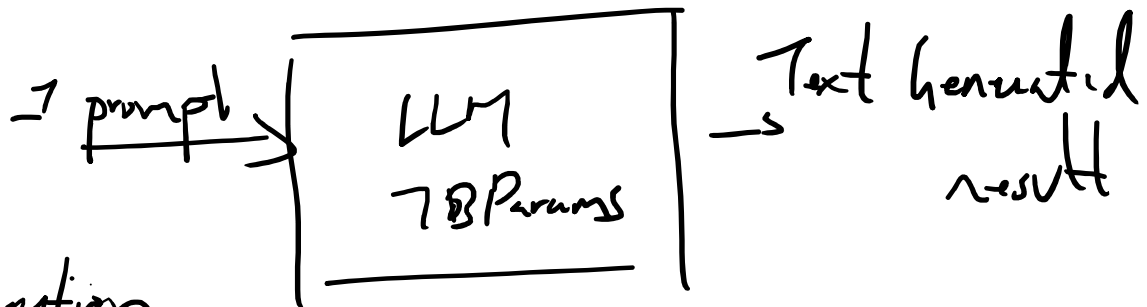
→ Colab
Local ~ Ollama

Vistara AI v Paid

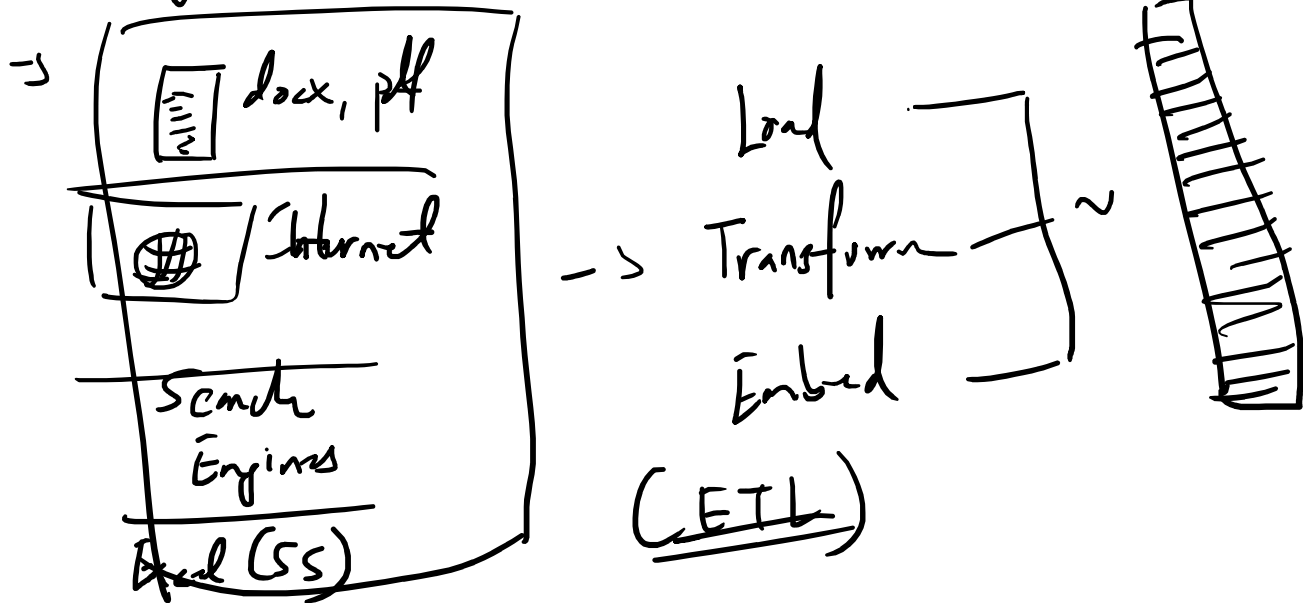
Local ~ ollama

→ RAH: Retrieval Augmented Generation

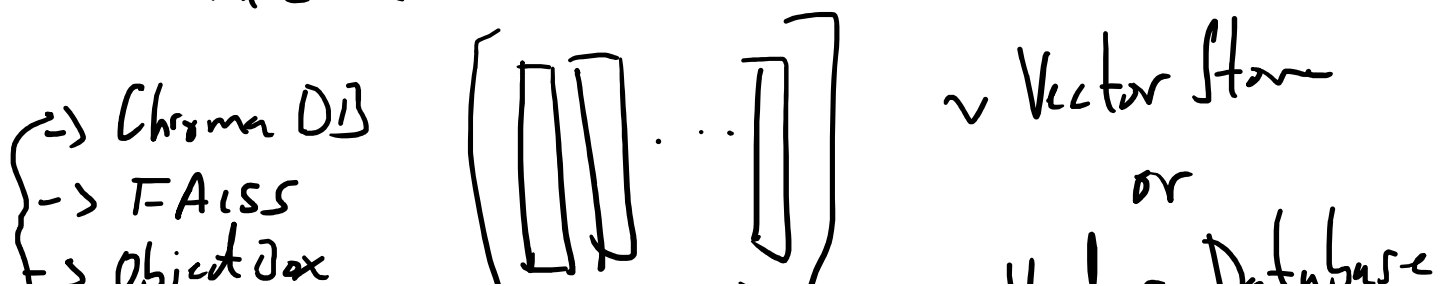
- Information is sparse



Data Ingestion



→ Load, Transform & Embed converts data into vectors.



-> H-ALSS
-> Object Box
-> Cassandra
(ASTRA)



Vector Database

Map & Reduce

-> Similarity Search

task: Create a query engine for paper.